

SaaS Leaders Must Trade Margin to Win the AI Race

10 March 2026 - ID G00846035 - 8 min read

By: Martin Prendergast, Mark Wesker

Initiatives: Technologies and Markets; Drive Next-Gen Business Models Through AI

Agentic AI is disrupting SaaS business models. For the SaaS CEO, the price of entry into the agentic AI club is lower margins. Many won't survive the transition. Those that do will deliver greater value to customers and have stronger traction and higher valuations.

Insights at a Glance

The move to agentic AI is reshaping software economics, making traditional SaaS gross margins averaging 85% or higher unrealistic. CEOs transitioning to agentic AI-led business models are advised to accept lowering their targeted gross margins to 40% to 60% to achieve traction and defensibility through deep workflow integration. Focus must shift to measuring agentic gross margin, net revenue retention (NRR), and lifetime value (LTV).

Key Insights and Data Points:

- AI-driven disruption is leading to lower valuations for SaaS companies and tougher fundraising for those not adapting to AI-first models, as investor expectations shift rapidly.
- Token costs are now a core part of COGS. Higher gross margins for agentic AI offerings often signal low usage and therefore low value and weak competitive positioning.
- Hybrid architectures and intelligent routing can reduce inference costs by 60% and mitigate third-party pricing risk. ¹
- Seat-based pricing is misaligned with AI-driven usage; consumption-based or hybrid models better align revenue with agent activity and protect against cost unpredictability.
- High NRR and LTV indicate strong adoption and customer value.

Recommendations:

- Position margin compression to 40% to 60% as the cost of building sticky, integrated AI solutions.
- Replace traditional seat-based pricing with consumption or hybrid pricing models to protect margins by aligning revenue with usage and manage cost volatility.
- Further reduce COGS with model-agnostic hybrid foundational model architectures while avoiding single-vendor dependency.

Strategic Planning Assumption

Through 2030, agentic AI companies will have 50% higher NRR than SaaS companies.

Impact Brief

For decades, SaaS valuation was driven by 80%+ gross margins. Today, tech CEOs adhering to this same benchmark are facing competitive pressure and potential down-rounds with investors punishing those that don't adapt to the new economic realities driven by AI. Change is happening and is driven by AI's ability to orchestrate workflows, posing risk to traditional SaaS models. The AI natives who are driving this disruption are scaling to \$100M ARR with initial gross margins as low as 25%, prioritizing market capture over margin efficiency.

Tech CEOs are facing a shrinking window to act. This is reflected in investor sentiment with SaaS stock performance, with the early 2026 sell-off in the sector seeing the forward P/E collapse from roughly 35x at the end of 2025 to around 20x. The impact is felt even more keenly in private SaaS companies which are typically valued at a 30% to 50% discount to their public peers due to the inherent illiquidity and execution risk. Raising capital without the AI roadmap is now harder than ever.

Late entrants to the AI race will face a consolidated market with top players in each sector beginning to dominate market share. But in acting, tech CEOs must be careful not to just add "AI wrappers" — commoditized tools with high margins but low levels of adoption and utilization, weak barriers to entry and high churn. And they must manage the impact of margin compression on working capital and cash flow, with 60% of growth-stage companies that fail to navigate the transition of their product and operational models to AI, being displaced by AI-native competitors by 2030. ²

Actions

Tech CEOs of SaaS companies evolving their strategy to win the agentic AI race must accept lower margins as the price. They must convince stakeholders that margin compression is in fact a strategic necessity, not a sign of a weak business model. It is the investment required to move from a commoditized software wrapper to a defensible, integrated system of action.

For a SaaS company to transition to being a true AI company, there is a fundamental shift in the operating model. AI evolves from a “feature” added to a product into the foundational layer that drives how the company creates value. Tech CEOs who attempt to preserve >80% gross margins are still thinking in terms of AI as a product feature and will likely fail to build the deep workflow integrations necessary to survive the coming “churn wave,” with net revenue retention (NRR) struggling to stay above 100% compared to the AI-first companies where NRR will exceed 150%.

As tech CEOs move their companies toward an AI-first model, they must mitigate some of the margin erosion through repositioning pricing and optimizing the AI foundational models linked to COGS. Over time, CEOs must also focus on using AI internally to improve efficiency and bring down labor costs.

To win the AI race, margin compression is a strategic necessity, not a sign of failure.

- **Position margin compression as the cost of workflow ownership:** The shift from “systems of record” to agentic “systems of action” naturally compresses gross margins from 85% to a 40% to 60% “sweet spot,” as the higher inference costs of autonomous execution become the necessary price for building deep, defensible workflow lock-in.
- **Protect margins: Shift pricing to consumption/hybrid models:** Replace legacy subscriptions with consumption-based or hybrid pricing to capture the variable costs of agentic AI and ensure revenue scales with the actual work performed by agents rather than human-login activity.
- **Optimize margins through model agnosticism:** Protect costs from AI foundational providers eroding margin, and then optimize for better commercial outcomes by implementing intelligent model routing. Direct routine tasks to cheaper, open-source models while reserving expensive foundational models only for complex reasoning.

How to Execute

The winners of the agentic AI race will be those who tailor offerings to industry-specific needs, successfully managing the margin compression to build unassailable moats.

Focus on the following key areas:

Position Margin Compression as the Cost of Workflow Ownership

Traditional SaaS margins were protected by low variable costs and therefore low COGS. In the agentic era, end users are moving from “systems of record” (storing data) to “systems of action” (executing work). With inference costs as the new COGS, a target margin of 40% to 60%, linked with high utilization, indicates your agents are performing essential, high-volume work that creates customer “stickiness.”

- **Audit your “moat”:** If your gross margins are >80%, audit your credit consumption. Low consumption is a “red flag” that your AI is an optional “toy” rather than a core workflow driver.
- **Establish financial guardrails:** Set a hard gross margin floor (e.g., 40% to 45%) and require CEO approval for any contract below this threshold. This prevents disciplined investment from becoming a “race to the bottom.” Also, the shift to greater usage incurs the Jevons paradox: Even as token costs are forecast to fall by 95% through 2030,³ as AI becomes cheaper, usage explodes, keeping your total costs high. CEOs must manage the associated cost and cash flow impact carefully. CEOs must also look at using AI internally to reduce GTM and services costs to further protect gross margins.

Protect Margins: Shift Pricing to Consumption/Hybrid Pricing

In an agentic environment, usage is decoupled from human activity; a single autonomous agent can perform thousands of operations, easily eclipsing the value of a seat-based subscription and drive margins into the red.

Pricing must align with work. To protect profitability and ensure your pricing scales with the value delivered, transition to one of the following frameworks:

- **Consumption-based pricing:** Link payment with the resources customers use such as tokens, API calls, or completed tasks, rather than a fixed, flat-rate subscription fee. This approach aligns COGS directly with usage, protecting gross margins by ensuring that as agentic activity and inference expenses scale, revenue scales in tandem.

- **Hybrid pricing:** Combines a predictable base subscription fee for platform access with a variable “usage tail” that charges for actual consumption such as credits, or tokens. This model provides stable cash flow while allowing revenue to scale automatically as the agent’s autonomous activity increases. By anchoring the baseline to platform availability and the variable portion to the “system of action” throughput, tech CEOs can balance revenue predictability with the ability to capture additional value from high-utility agentic workflows.

Optimize Margins Through Foundational Model Agnosticism

Relying on a single foundational model leaves your agentic gross margin vulnerable to a third party’s pricing. Successful midmarket players use “intelligent routing” to arbitrage costs and improve agentic gross margins.

- **Route by complexity:** Direct 80% of routine tasks to cheaper, open-source models (like Llama 3) or optimized smaller models, reserving expensive foundational models only for complex reasoning tasks. This can restore significant margin points.
- **Build a hybrid deterministic architecture:** Use deterministic rules for known, repetitive queries (zero inference cost) and only trigger LLMs for ambiguous tasks. This can reduce inference costs by up to 60%. (Learn more: [Comparing Deterministic and Agentic Options for Automating Business Processes](#)).
- **Own the evaluation:** Build “private evals” based on your proprietary data. This allows you to swap underlying models confidently without breaking the user experience, ensuring you always access the lowest-cost compute for the required performance.

Success Measures

Tech CEOs must focus on the following metrics to measure success:

- **Agentic gross margin:** Unlike SaaS gross margin, this is calculated by subtracting the specific inference costs (tokens, GPU time, API latency overhead) from the revenue generated by that specific task or workflow. It discounts COGS not directly linked to inference and will indicate the success of foundational model agnosticism.
- **Net revenue retention (NRR):** Target >150%. NRR captures whether customers are scaling their autonomous workflows beyond initial use cases, signaling that the agents are delivering the tangible value to justify expanded consumption, and there is strong “stickiness” in the agentic solution.

- **Lifetime value (LTV):** LTV measures the long-term compounding value and indicates that the product has achieved deep operational “stickiness” within the client’s stack. As agents successfully automate increasingly complex tasks, the resulting expansion in usage and NRR drives a higher LTV.

Evidence

¹ [Optimize AI Cost and Reliability Using AI Gateways and Model Routers.](#)

² [AI Vendor Race: The AI Investment Framework: How Growth Tech CEOs Win.](#)

³ [Token Costs Will Fall by 95% Through 2030.](#)

Recommended by the Authors

Some documents may not be available as part of your current Gartner subscription.

[Tech CEOs Must Outmaneuver AI Giants by Focusing on Microindustry Outcomes](#)

© 2026 Gartner, Inc. and/or its affiliates. All rights reserved. Gartner is a registered trademark of Gartner, Inc. and its affiliates. This publication may not be reproduced or distributed in any form without Gartner's prior written permission. It consists of the opinions of Gartner's Business and Technology Insights Organization, which should not be construed as statements of fact. While the information contained in this publication has been obtained from sources believed to be reliable, Gartner disclaims all warranties as to the accuracy, completeness or adequacy of such information. Although Gartner insights may address legal and financial issues, Gartner does not provide legal or investment advice and its insights should not be construed or used as such. Your access and use of this publication are governed by [Gartner's Usage Policy](#). Gartner prides itself on its reputation for independence and objectivity. Its insights is produced independently by its Business and Technology Insights Organization without input or influence from any third party. For further information, see "[Guiding Principles on Independence and Objectivity](#)." Gartner insights may not be used as input into or for the training or development of generative artificial intelligence, machine learning, algorithms, software, or related technologies.